AI CODE EVALUATION PLATFORM

Two Sum - Evaluation report

Submission label	Buggy example
Submission ID	544e950e-6197-4289-8559-d83324f0fde9
Challenge	Two Sum
Evaluated (UTC)	2026-10-17T18:27:40.558882+00:00
Source label	trusted-example
Mode	trusted
Test pass rate	3/5 (60.00%)
Completed-case time	4.421 ms
Ruff	completed; 0 findings
Reference verified	True
Runtime	CPython 3.12.14
Candidate SHA-256	f580db2ba43bf9e93e32546cfa4ca680 39a34477eed83c01fd88812ce583a75f

3/5 deterministic tests passed. Test pass rate is not code coverage or proof of general correctness. Review the reproducible failures below.

Deterministic test pass rate is not code coverage or proof of correctness. Timings are single measurements, include interpreter work and exclude container startup. Timed-out or resource-terminated cases have no completed duration. RSS includes the interpreter (Linux KiB). Hidden edge cases are revealed after evaluation for reproducible debugging.

basic - passed

hidden: False

input: {"nums": [2, 7, 11], "target": 9}

expected: [0, 1]

actual: [0, 1]

error: None

duration_ms: 0.921244

reference_duration_ms: 0.901825

peak_rss_kib: 42496

no-reuse - wrong_output

hidden: True

input: {"nums": [3], "target": 6}

expected: []

actual: [0, 0]

error: Returned value differs from the challenge contract

duration_ms: 0.806291

reference_duration_ms: 1.087583

peak_rss_kib: 42496

repeated - wrong_output

hidden: True

input: {"nums": [3, 3], "target": 6}

expected: [0, 1]

actual: [0, 0]

error: Returned value differs from the challenge contract

duration_ms: 0.977879

reference_duration_ms: 0.913202

peak_rss_kib: 42496

empty - passed

hidden: True

input: {"nums": [], "target": 0}

expected: []

actual: []

error: None

duration_ms: 0.816106

reference_duration_ms: 0.794683

peak_rss_kib: 42496

negative - passed

hidden: True

input: {"nums": [-4, 2, 5], "target": 1}

expected: [0, 2]

actual: [0, 2]

error: None

duration_ms: 0.899091

reference_duration_ms: 1.034243

peak_rss_kib: 42496

Static analysis

[]

Candidate source

```
def solve(nums, target):
    seen = {}
    for i, value in enumerate(nums):
        seen[value] = i
        if target - value in seen:
            return [seen[target - value], i]
    return []
```

Execution limits

{
"wall_seconds_per_case": 3,
"cpu_seconds_per_case": 1,
"memory_mib": 128,
"output_bytes": 65536,
"docker_pids": 16,
"network": "none",
"scope": "per test; RSS includes interpreter; Linux KiB; time excludes container startup"
}